

✓ Frame-Level Speech Recognition

✓ EXPERIMENT

- Build LA-DNN
 - A powerful model for speech recognition and good candidate to replace FF-DNN
 - Helps to remove the vanishing gradient problems in deep network
 - <https://ieeexplore.ieee.org/document/7472646>
- No Augmentation
- No hyper-parameter tuning however used our past knowledge to set hyper-parameters
- Advanced weight initialization, Label smoothing in loss function, gradient clipping and scheduler are used.

```
1 !pip install torch-summary --quiet
```

```
1 import torch
2 import numpy as np
3 import sklearn
4 import gc
5 import zipfile
6 import pandas as pd
7 from tqdm.auto import tqdm
8 import os
9 import datetime
10 import torchsummary
11 device = 'cuda' if torch.cuda.is_available() else 'cpu'
12 print("Device: ", device)
```

⇒ Device: cuda

```
1 # ### If you are using colab, you can import google drive
2 from google.colab import drive
3 drive.mount('/content/drive')
```

```
1 !pip install --upgrade --force-reinstall --no-deps kaggle==
2 !mkdir -p /root/.kaggle
```

⇒ Collecting kaggle==1.5.8
 Downloading kaggle-1.5.8.tar.gz (59 kB)
 59.2/59.2 kB 1.8 MB/s eta 0:00:

```
Preparing metadata (setup.py) ... done
Building wheels for collected packages: kaggle
  Building wheel for kaggle (setup.py) ... done
  Created wheel for kaggle: filename=kaggle-1.5.8-py3-none-any.whl size=73249
  Stored in directory: /root/.cache/pip/wheels/b5/23/bd/d33cbf399584fa44fa049
Successfully built kaggle
Installing collected packages: kaggle
  Attempting uninstall: kaggle
    Found existing installation: kaggle 1.7.4.2
    Uninstalling kaggle-1.7.4.2:
      Successfully uninstalled kaggle-1.7.4.2
Successfully installed kaggle-1.5.8
```

```
1 with open("/root/.kaggle/kaggle.json", "w+") as f:
2     f.write('{"username":"dineshbuswala","key":"a67fefaecac
3
4 !chmod 600 /root/.kaggle/kaggle.json
```

```
1 # commands to download data from kaggle
2
3 !kaggle competitions download -c 11785-hw1p2-f23 --force
4 !mkdir -p '/content/data'
5
6 # !unzip -qo /content/11785-hw1p2-f23.zip -d '/content/data'
```

```
Downloaded 11785-hwlp2-f23.zip to /kaggle/working  
100%|██████████| 3.99G/3.99G [00:24<00:00, 219MB/s]  
100%|██████████| 3.99G/3.99G [00:24<00:00, 174MB/s]
```

```
1 !unzip -go /kaggle/working/11785-hw1p2-f23.zip -d '/content'
```

```
1 # ! rm -r /content/data/11-785-f23-hw1p2/test-clean
2 # ! rm -r /content/11785-hw1p2-f23.zip
```

```
1 ! ls /content/data/11-785-f23-hw1p2/train-clean-100/mfcc/
```

28539

```
1 ### configuration variables
2 EPOCHS = 4
3 BATCH_SIZE = 2048 * 2
4 LEFT_CONTEXT = 7
5 RIGHT_CONTEXT = 7
6 INITIAL_LEARNING_RATE = 1e-3
7 L2_PENALTY = 1e-5
8 STEP_SIZE = 2
```

```

9 GAMMA = 0.1
10 BASE_DIRECTORY = '/content/data/11-785-f23
11 TRAINING_DATA = BASE_DIRECTORY + 'train-c
12 EVALUATION_DATA = BASE_DIRECTORY + 'dev-cle
13 PHONEMES = ['[SIL]', 'AA', 'AE'
14             'B', 'CH', 'D',
15             'F', 'G', 'HH',
16             'L', 'M', 'N',
17             'R', 'S', 'SH',
18             'V', 'W', 'Y',
19 PHONEMES_TO_INDEX = {phoneme: idx for idx, p
20 NUMBER_OF_NEURONS = [2048, 2048, 1024, 1024,
21 MODEL_DIR = "/content"
22 CLIP_VALUE = 1.0
23 LABEL_SMOOTHING = 0.01

```

```

1 class AudioDataset(torch.utils.data.Dataset):
2     def __init__(self, root, phonemes = PHONEMES_TO_INDEX,
3                 left_context = LEFT_CONTEXT,
4                 right_context = RIGHT_CONTEXT):
5
6         self.left_context = left_context
7         self.right_context = right_context
8         self.phonemes_mapping = phonemes
9
10        self.mfcc_dir = os.path.join(root, 'mfcc')
11        self.transcript_dir = os.path.join(root, 'transcri
12
13        # List and sort mfcc and transcript files
14        mfcc_names = sorted(os.listdir(self.mfcc_dir))
15        transcript_names = sorted(os.listdir(self.transcri
16
17        # Sanity check
18        assert len(mfcc_names) == len(transcript_names), "
19
20        total_frames = 0
21
22        for i in range(len(mfcc_names)):
23            mfcc_path = os.path.join(self.mfcc_dir, mfcc_n
24
25            mfcc = np.load(mfcc_path,
26                          allow_pickle = False,
27                          mmap_mode='r')
28            total_frames += mfcc.shape[0]
29            del mfcc

```

```

30
31     sample_mfcc = np.load(os.path.join(self.mfcc_dir,
32     self.mfcc_dim = sample_mfcc.shape[1]
33     ### Right Padding is added here automatically
34     self.mfccs = np.zeros((total_frames + right_context
35     self.transcripts = [None] * (total_frames + right_
36     ### Release memory
37     del sample_mfcc, total_frames
38     gc.collect()
39
40     index = 0
41     for i in range(len(mfcc_names)):
42         mfcc = np.load(os.path.join(self.mfcc_dir, mfc
43             allow_pickle = False,
44             mmap_mode = 'r')
45
46         mfcc = (mfcc - mfcc.mean(axis = 1, keepdims =
47
48         transcript = np.load(os.path.join(self.transcr
49             allow_pickle = False,
50             mmap_mode='r'))[1:-1]
51
52         self.mfccs[index: index + mfcc.shape[0]] = mfc
53         self.transcripts[index: index + len(transcript
54         index += mfcc.shape[0]
55         del mfcc, transcript
56
57         if i % 1000 == 0:
58             gc.collect() ### Release memory
59
60     # Save original dataset length (before adding righ
61     self.length = len(self.mfccs) - right_context
62
63     # Map transcript phonemes to indices
64     self.transcripts = [self.phonemes_mapping.get(p, -
65
66     ### Release memory
67     del mfcc_names, transcript_names
68     gc.collect()
69
70     def __len__(self):
71         return self.length
72
73     def __getitem__(self, ind):
74         if ind < self.left_context: # index is less than

```

```

75         # zeros need to prepend with it
76         padding = []
77         for _ in range(self.left_context - ind):
78             padding.append(np.zeros((1, self.mfcc_dim)
79
80         for i in range(ind):
81             padding.append(self.mfccs[i][None, :])
82
83         # Include current and right context frames
84         current = self.mfccs[ind][None, :] # Ensure s
85         right_context = [self.mfccs[i][None, :] for i
86         frames = np.concatenate(padding + [current] +
87
88     else: # when index is greater than or equal to le
89         left_context_frames = self.mfccs[ind - self.le
90         current = self.mfccs[ind][None, :] # Ensure s
91         right_context_frames = self.mfccs[ind + 1: ind
92
93         frames = np.concatenate([left_context_frames,
94
95         frames = frames.flatten() # Flatten to get 1D dat
96         frames = torch.FloatTensor(frames) # Convert to t
97         phonemes = torch.tensor(self.transcripts[ind], dt
98
99         return frames, phonemes
100

```

```

1 # Create a dataset object using the AudioDataset class for
2 train_data = AudioDataset(TRAINING_DATA)
3
4 # Create a dataset object using the AudioDataset class for
5 val_data = AudioDataset(EVALUATION_DATA)
6

```

```

1 train_loader = torch.utils.data.DataLoader(
2     dataset      = train_data,
3     num_workers  = 2,
4     batch_size   = BATCH_SIZE,
5     pin_memory   = True,
6     shuffle      = True
7 )
8
9 val_loader = torch.utils.data.DataLoader(
10     dataset      = val_data,
11     num_workers  = 1,

```

```

12     batch_size = BATCH_SIZE,
13     pin_memory = True,
14     shuffle     = False
15 )
16
17 print("Train dataset samples = {}, batches = {}".format(tra
18 print("Validation dataset samples = {}, batches = {}".forma

```

→ Train dataset samples = 36091157, batches = 8812
 Validation dataset samples = 1928204, batches = 471

```

1 # # Testing code to check if data loaders are working as ex
2 # total_batches = len(train_loader)
3
4 # for i, (frames, phoneme) in enumerate(train_loader):
5 #     if i < 10 or i >= total_batches - 10:
6 #         print(f"Batch {i + 1}/{total_batches}")
7 #         print("Frames:", frames)
8 #         print("Phoneme:", phoneme)
9 #         print("-" * 50)
10

```

```

1 class LADNNLayer(torch.nn.Module):
2     def __init__(self, in_features, out_features, rank = -1
3         super(LADNNLayer, self).__init__()
4
5         self.in_features = in_features
6         self.out_features = out_features
7         self.activation = torch.nn.ReLU()
8         if rank == -1:
9             self.intermediate = in_features // 2
10        else:
11            self.intermediate = rank
12        self.U = torch.nn.Linear(self.in_features, self.int
13        self.V = torch.nn.Linear(self.intermediate, self.ou
14        self.T = torch.nn.Linear(self.in_features, self.out
15
16        # Apply weight initialization
17        self._initialize_weights()
18
19        def _initialize_weights(self):
20            print('Initialization of weights using Kaiming')
21            for m in [self.U] + [self.V] + [self.T]:
22                if isinstance(m, torch.nn.Linear):
23                    # Kaiming initialization for weights

```

```

24         torch.nn.init.kaiming_uniform_(m.weight, no
25
26     def forward(self, x):
27         t_x = self.T(x)
28         u_x = self.U(x)
29         v_x = self.V(self.activation(u_x))
30         return v_x + t_x
31
32
33
34 class Network(torch.nn.Module):
35     def __init__(self, input_size):
36         super(Network, self).__init__()
37
38         # Neurons in each layer: input -> hidden(s) -> outp
39         self.neurons = [input_size] + NUMBER_OF_NEURONS + [
40
41         layers = []
42         for in_features, out_features in zip(self.neurons[:
43             layers.append(LADNNLayer(in_features, out_featu
44
45         # Final layer
46         layers.append(LADNNLayer(self.neurons[-2], self.neu
47
48         # Combine all into a sequential model
49         self.model = torch.nn.Sequential(*layers)
50
51         # Apply weight initialization
52         self._initialize_weights()
53
54     def _initialize_weights(self):
55         print('Initialization of weights using Kaiming')
56         for m in self.model:
57             if isinstance(m, torch.nn.Linear):
58                 # Kaiming initialization for weights
59                 torch.nn.init.kaiming_uniform_(m.weight, no
60
61
62     def forward(self, x):
63         out = self.model(x)
64         return out
65

```

```

1 INPUT_SIZE = (LEFT_CONTEXT + RIGHT_CONTEXT + 1) * 28
2 model = Network(INPUT_SIZE).to(device)

```

3 # Pass the input size as a tuple, without the batch dimensi

4 torchsummary.summary(model, (INPUT_SIZE,))

```

=> Initialization of weights using Kaiming
Initialization of weights using Kaiming
Initialization of weights using Kaiming
Initialization of weights using Kaiming
Initialization of weights using Kaiming
Initialization of weights using Kaiming
Initialization of weights using Kaiming
Initialization of weights using Kaiming
Initialization of weights using Kaiming

```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 2048]	862,208
Linear-2	[-1, 210]	88,410
ReLU-3	[-1, 210]	0
Linear-4	[-1, 2048]	432,128
LADNNLayer-5	[-1, 2048]	0
Linear-6	[-1, 2048]	4,196,352
Linear-7	[-1, 1024]	2,098,176
ReLU-8	[-1, 1024]	0
Linear-9	[-1, 2048]	2,099,200
LADNNLayer-10	[-1, 2048]	0
Linear-11	[-1, 1024]	2,098,176
Linear-12	[-1, 1024]	2,098,176
ReLU-13	[-1, 1024]	0
Linear-14	[-1, 1024]	1,049,600
LADNNLayer-15	[-1, 1024]	0
Linear-16	[-1, 1024]	1,049,600
Linear-17	[-1, 512]	524,800
ReLU-18	[-1, 512]	0
Linear-19	[-1, 1024]	525,312
LADNNLayer-20	[-1, 1024]	0
Linear-21	[-1, 512]	524,800
Linear-22	[-1, 512]	524,800
ReLU-23	[-1, 512]	0
Linear-24	[-1, 512]	262,656
LADNNLayer-25	[-1, 512]	0
Linear-26	[-1, 256]	131,328
Linear-27	[-1, 256]	131,328
ReLU-28	[-1, 256]	0
Linear-29	[-1, 256]	65,792
LADNNLayer-30	[-1, 256]	0
Linear-31	[-1, 256]	65,792
Linear-32	[-1, 128]	32,896
ReLU-33	[-1, 128]	0
Linear-34	[-1, 256]	33,024
LADNNLayer-35	[-1, 256]	0
Linear-36	[-1, 40]	10,280
Linear-37	[-1, 128]	32,896
ReLU-38	[-1, 128]	0
Linear-39	[-1, 40]	5,160
LADNNLayer-40	[-1, 40]	0

Total params: 18,942,890

Trainable params: 18,942,890

Non-trainable params: 0

Train size (MB): 0.00

```

1 criterion = torch.nn.CrossEntropyLoss(label_smoothing = LAB
2 # We use CE because the task is multi-class classification
3
4 optimizer = torch.optim.Adam(model.parameters(),
5                               lr = INITIAL_LEARNING_RATE,
6                               weight_decay = L2_PENALTY) #De
7 scheduler = torch.optim.lr_scheduler.StepLR(
8     optimizer, step_size = STEP_SIZE, gamma = GAMMA
9 )
10 # Refer - https://pytorch.org/docs/stable/notes/amp_example

```

```

1 torch.cuda.empty_cache()
2 gc.collect()

```

 101

```

1 def train(model, dataloader, optimizer, criterion):
2
3     model.train()
4     tloss, tacc = 0, 0 # Monitoring loss and accuracy
5     batch_bar = tqdm(total=len(train_loader), dynamic_ncol=
6
7     for i, (frames, phonemes) in enumerate(dataloader):
8
9         ### Initialize Gradients
10        optimizer.zero_grad()
11
12        ### Move Data to Device (Ideally GPU)
13        frames = frames.to(device)
14        phonemes = phonemes.to(device)
15
16        ### Forward Propagation
17        logits = model(frames)
18
19        ### Loss Calculation
20        loss = criterion(logits, phonemes)
21
22        ### Backward Propagation
23        loss.backward()
24
25        ### Clip gradients
26        torch.nn.utils.clip_grad_norm_(model.parameters(),
27
28        ### Gradient Descent

```

```

29         optimizer.step()
30
31         tloss    += loss.item()
32         tacc     += torch.sum(torch.argmax(logits, dim= 1) ==
33                               target).item()
34         batch_bar.set_postfix(loss="{:.04f}".format(float(tloss/acc)),
35                               acc="{:.04f}%".format(float(tacc/tloss)))
36         batch_bar.update()
37
38         ### Release memory
39         del frames, phonemes, logits
40         torch.cuda.empty_cache()
41
42     batch_bar.close()
43     tloss    /= len(train_loader)
44     tacc     /= len(train_loader)
45
46     return tloss, tacc

```

```

1 def eval(model, dataloader):
2
3     model.eval() # set model in evaluation mode
4     vloss, vacc = 0, 0 # Monitoring loss and accuracy
5     batch_bar    = tqdm(total=len(val_loader), dynamic_ncols=100)
6
7     for i, (frames, phonemes) in enumerate(dataloader):
8
9         ### Move data to device (ideally GPU)
10        frames      = frames.to(device)
11        phonemes    = phonemes.to(device)
12
13        # makes sure that there are no gradients computed a
14        with torch.inference_mode():
15            ### Forward Propagation
16            logits  = model(frames)
17            ### Loss Calculation
18            loss    = criterion(logits, phonemes)
19
20            vloss   += loss.item()
21            vacc    += torch.sum(torch.argmax(logits, dim= 1) ==
22                                target).item()
23
24            batch_bar.set_postfix(loss="{:.04f}".format(float(vloss/vacc)),
25                                  acc="{:.04f}%".format(float(vacc/vloss)))
26            batch_bar.update()

```

```
27
28     ### Release memory
29     del frames, phonemes, logits
30     torch.cuda.empty_cache()
31
32     batch_bar.close()
33     vloss /= len(val_loader)
34     vacc /= len(val_loader)
35
36     return vloss, vacc

1 best_model_path = os.path.join(MODEL_DIR, "best_model.pt")
2 best_acc = -np.inf
3 for epoch in range(EPOCHS):
4     ### clean up memory before computation
5     torch.cuda.empty_cache()
6     gc.collect()
7
8
9     print(f"\nEpoch {epoch + 1}/{EPOCHS}")
10
11     curr_lr = float(optimizer.param_groups[0]['lr'])
12     train_loss, train_acc = train(model, train_loader, opti
13     val_loss, val_acc = eval(model, val_loader)
14
15     print(f"\tTrain Acc: {train_acc*100:.2f}%\tTrain Loss:
16     print(f"\tVal    Acc: {val_acc*100:.2f}%\tVal    Loss: {v
17
18     # Save model at every epoch
19     epoch_model_path = os.path.join(MODEL_DIR, f"model_at_e
20     torch.save(model.state_dict(), epoch_model_path)
21
22     # Save best model
23     if val_acc > best_acc:
24         best_acc = val_acc
25         torch.save(model.state_dict(), best_model_path)
26         print(f"Updated Best Model at: {best_model_path}")
27
28     ### take step in adjusting the learning rate
29     scheduler.step()
30
31
```



```
Epoch 1/4
Train: 0%|          | 0/8812 [00:00<?, ?it/s]
Val: 0%|          | 0/471 [00:00<?, ?it/s]
      Train Acc: 67.04%      Train Loss: 13.0031      LR: 0.0010000
      Val  Acc: 69.21%      Val  Loss: 1.0357
Updated Best Model at: /content/best_model.pt
```

```
Epoch 2/4
Train: 0%|          | 0/8812 [00:00<?, ?it/s]
Val: 0%|          | 0/471 [00:00<?, ?it/s]
      Train Acc: 72.84%      Train Loss: 0.9069      LR: 0.0010000
      Val  Acc: 70.96%      Val  Loss: 0.9690
Updated Best Model at: /content/best_model.pt
```

```
Epoch 3/4
Train: 0%|          | 0/8812 [00:00<?, ?it/s]
Val: 0%|          | 0/471 [00:00<?, ?it/s]
      Train Acc: 77.41%      Train Loss: 0.7553      LR: 0.0001000
      Val  Acc: 74.43%      Val  Loss: 0.8521
Updated Best Model at: /content/best_model.pt
```

```
Epoch 4/4
Train: 0%|          | 0/8812 [00:00<?, ?it/s]
Val: 0%|          | 0/471 [00:00<?, ?it/s]
      Train Acc: 78.37%      Train Loss: 0.7237      LR: 0.0001000
      Val  Acc: 74.58%      Val  Loss: 0.8479
Updated Best Model at: /content/best_model.pt
```

1

2 # Load best model after training

3 model.load_state_dict(torch.load('/content/best_model.pt'))

4

⇒ <All keys matched successfully>

```
1 from sklearn.metrics import classification_report, confusion_matrix
2 model.eval() # Set model in evaluation mode
3 predicted = []
4 groundtruth = []
5
6 for frames, phonemes in val_loader:
7
8     # Move data to device
9     frames = frames.to(device)
10    phonemes = phonemes.to(device)
11
12    # Disable gradient calculation
13    with torch.inference_mode():
14        logits = model(frames)
15
16    predict = torch.argmax(logits, dim = 1)
17
```

```

18 # Detach and move to CPU for evaluation
19 predicted.extend(predict.detach().cpu().tolist())
20 groundtruth.extend(phonemes.detach().cpu().tolist())
21
22 # Release memory
23 del frames, phonemes, logits, predict
24 torch.cuda.empty_cache()
25
26 # Print classification report
27 print(classification_report(
28     groundtruth,
29     predicted,
30     target_names = PHONEMES # Skipping SOS and EOS tokens
31 ))
32

```



	precision	recall	f1-score	support
[SIL]	0.94	0.95	0.94	319908
AA	0.59	0.59	0.59	29688
AE	0.65	0.66	0.66	49298
AH	0.61	0.64	0.63	123734
AO	0.66	0.64	0.65	29340
AW	0.72	0.63	0.67	20274
AY	0.79	0.83	0.81	49332
B	0.69	0.68	0.68	23607
CH	0.65	0.62	0.63	12644
D	0.65	0.58	0.61	62763
DH	0.70	0.68	0.69	37100
EH	0.62	0.59	0.60	47112
ER	0.69	0.68	0.69	54928
EY	0.75	0.76	0.76	36184
F	0.72	0.78	0.75	37562
G	0.74	0.69	0.71	13541
HH	0.75	0.70	0.72	34813
IH	0.63	0.59	0.61	74887
IY	0.77	0.79	0.78	70861
JH	0.69	0.61	0.65	8730
K	0.77	0.80	0.79	47016
L	0.73	0.82	0.77	65902
M	0.77	0.78	0.78	44728
N	0.74	0.77	0.75	94541
NG	0.73	0.69	0.71	19327
OW	0.71	0.62	0.66	30755
OY	0.77	0.54	0.63	3861
P	0.75	0.70	0.72	34131
R	0.69	0.75	0.72	62686
S	0.80	0.82	0.81	101184
SH	0.81	0.81	0.81	17628
T	0.69	0.68	0.68	97390
TH	0.44	0.40	0.42	9247
UH	0.64	0.45	0.53	6286
UW	0.73	0.67	0.70	26691
V	0.71	0.63	0.67	27440
W	0.78	0.81	0.80	37697

Y	0.70	0.60	0.65	9669
Z	0.71	0.70	0.70	54850
ZH	0.71	0.55	0.62	869
accuracy			0.75	1928204
macro avg	0.71	0.68	0.69	1928204
weighted avg	0.74	0.75	0.74	1928204

```

1 import seaborn as sns
2 import matplotlib.pyplot as plt
3 sns.set_style("darkgrid")

```

```

1 # Compute confusion matrix
2 cm = confusion_matrix(groundtruth, predicted)
3
4 # Normalize confusion matrix by row (i.e., by true labels)
5 cm_normalized = cm.astype('float') / cm.sum(axis = 1, keepd
6
7 # Replace NaNs (from division by zero, if any row sum is 0)
8 cm_normalized = np.nan_to_num(cm_normalized)
9
10 # Plot
11 plt.figure(figsize=(15, 16))
12 sns.heatmap(cm_normalized,
13             annot=True,
14             fmt=".1f",
15             cmap="Greens",
16             xticklabels=PHONEMES,
17             yticklabels=PHONEMES,
18             cbar_kws={'label': 'Proportion'})
19
20 plt.title('Normalized Confusion Matrix', fontsize=16)
21 plt.xlabel('Predicted Label', fontsize=12)
22 plt.ylabel('True Label', fontsize=12)
23 plt.xticks(rotation=45, ha='right', fontsize=10)
24 plt.yticks(rotation=0, fontsize=10)
25 plt.tight_layout()
26 plt.show()
27

```





